

GeoCalib: Calibrating Geometric Consistency for Reliable 3D Scene Graph Relations

Anonymous submission

Abstract

3D Scene Graphs represent objects and relations in a form useful for spatial reasoning, but relation predictions can remain unreliable even when predicted predicates appear semantically plausible. We study a relation-level failure mode: semantic relation confidence is not necessarily calibrated to the 3D geometry of the same object pair, so a predictor can rank physically inconsistent relations highly. To expose and reduce this failure, we introduce a GeoCalib framework for geometry-checkable relation families. The framework factorizes predicate semantics T , predicate-independent same-pair geometry G , and source confidence Z ; it calibrates a constructed-target compatibility $C = P(y^{cal} = 1 \mid T, G)$ without using Z , then reports $F(Z, C)$ under a recall-violation protocol. Two soft fusion instantiations use a calibrated product and a scale-robust rank-average; pooled calibration, Reciprocal Rank Fusion, and hard rules serve as ablation, comparison, and diagnostic. On the full official validation split, the calibrated product improves VL-SAT R@50/R@100 from 0.9272/0.9635 to 0.9288/0.9683 while reducing Violation@100 from 0.0476 to 0.0333. On the main open-vocabulary relation-source case study, the Open3DSG full-validation recovery branch, the product improves R@100 from 0.5161 to 0.6047 while reducing Violation@100 from 0.1242 to 0.0341 under explicit selected-checkpoint and recovery-policy caveats. On a prospectively frozen SGFN target, both the product and pre-specified rank-average pass the joint R@100/verifier-Violation@100 gate, showing that the aggregate effect is not tied to one fusion formula. GT positive/counterfactual checks separate valid and invalid geometry with AUROC/AUPRC of 0.9772/0.9729, and controls rule out calibrator-only, distance-only, shuffled-geometry, and wrong-pair explanations. Here y^{cal} denotes train-dev GT positives and high-margin counterfactual negatives, not direct human validity. These results support a scoped relation-reliability claim for measured geometry-checkable 3D Scene Graph relations, not formula dominance or broad open-vocabulary graph generation.

Introduction

3D Scene Graphs turn reconstructed scenes into structured object-relation representations that can support spatial reasoning, embodied AI, scene alignment, visual grounding, and

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

downstream decision making (Wald et al. 2020; Sarkar et al. 2023; Xie, Pagani, and Stricker 2024; Liu et al. 2026). For these uses, relation prediction must be reliable in the physical scene, not only plausible at the category or language level. A graph edge such as `standing on`, `next to`, or `higher than` is useful only if it describes the actual subject-object pair in 3D.

We study a failure mode in which semantic relation predictors rank plausible relation labels that are geometrically inconsistent for the actual object pair. A relation can sound plausible for two object categories while being contradicted by their contact evidence, distance, or vertical arrangement in the reconstructed scene. This mismatch matters more as 3D Scene Graphs become interfaces for language-facing, open-vocabulary, and embodied systems (Koch et al. 2024b; Hou et al. 2025; Zhang et al. 2025; Wang et al. 2025; Saxena and Chiun 2025; Yu et al. 2025; Madhavaram et al. 2026).

The failure is not simply that existing 3DSSG methods ignore geometry. Prior work already uses edge features, point-cloud structure, spatial knowledge, multimodal semantics, and geometric fusion (Zhang et al. 2021; Feng et al. 2023; Wang et al. 2023; Koch et al. 2024a; Chen et al. 2024). The gap is more specific: semantic relation confidence is not necessarily a calibrated estimate of relation-level physical consistency. A high relation score can behave as a semantic plausibility score without indicating whether the same subject-object pair satisfies the relation in 3D. This makes the problem an AI reliability issue, not only a 3DSSG implementation detail: language-facing and embodied systems can inherit a confident symbolic relation that contradicts the underlying world state.

This motivates a GeoCalib framework. The framework treats an existing predictor as a relation source, standardizes its outputs into identity-preserving prediction rows, joins object-pair geometry evidence by subject and object instance identifiers, and separates predicate/family semantics T , predicate-independent same-pair geometry G , and source confidence Z . A compatibility factor $C = P(y^{cal} = 1 \mid T, G)$ is calibrated without Z ; source confidence enters only in the final ranking $F(Z, C)$. GeoCalib also separates the framework from any one fusion formula. A calibrated product preserves semantic-score magnitudes while penalizing family-conditioned geometry risk; a pre-specified rank-average combines semantic and geometry percentiles

without assuming shared source scales. Hard physical rules diagnose strict consistency, pooled calibration is an ablation, and Reciprocal Rank Fusion is a strong comparison. Figure 1 summarizes this path.

We evaluate the framework on geometry-checkable relation families: `support_contact`, `proximity`, and `relative_vertical`. These families are selected because their validity can be tested from reconstructed object-pair geometry, they cover contact, metric proximity, and vertical ordering, and they appear in both relation sources with a fixed exact-label denominator. Open3DSG is the main open-vocabulary relation source case study, and VL-SAT is a controlled reproduced closed-set anchor within the same geometry-checkable family scope (Wang et al. 2023; Koch et al. 2024b). Controls test whether the effect can be explained by calibrator-only ranking, a generic distance heuristic, shuffled geometry, or wrong-pair geometry. GT-positive and counterfactual checks test the geometry-validity signal independently from prediction re-ranking, while qualitative failure analysis shows how semantic plausibility diverges from physical consistency.

Our contributions are threefold:

1. We identify and formalize a relation-level reliability failure in 3D Scene Graph prediction: semantic relation confidence can rank plausible predicates that are physically inconsistent because it is not calibrated to object-pair geometry.
2. We introduce a GeoCalib framework that standardizes prediction rows, factorizes T , raw G , and Z , calibrates $C(T, G)$ under the leakage boundary $Z \notin C$, and exposes calibrated-product and scale-robust rank-average soft fusion instantiations while retaining pooled, RRF, and rule-verified comparisons.
3. We define a recall-violation evaluation protocol for geometry-checkable relation families, including exact-label $R@K$, Violation@K, GT-positive/counterfactual verifier evaluation, and geometry identity controls.

We validate these contributions with Docker-reproducible evidence on VL-SAT, Open3DSG, and a prospectively frozen SGFN target, including nontriviality controls, family decomposition, failure analysis, and explicit reproduction caveats. The claim is intentionally scoped: GeoCalib measures and improves relation reliability for geometry-checkable 3DSSG families. It does not claim general open-vocabulary 3D Scene Graph generation improvement, arbitrary-source generality, formula dominance, or guaranteed physical correctness.

Related Work

3D Scene Graph relation prediction. 3D Scene Graphs organize indoor scenes as object-relation structures grounded in 3D space (Armeni et al. 2019). The 3RScan/3DSSG setting established predicate recall for 3D relation prediction (Wald et al. 2020), while incremental and online systems study how scene graphs can be constructed from RGB-D streams or spatial perception stacks (Wu et al. 2021; Hughes, Chang, and Carbone 2022). Recent predictors improve closed-set predicates using edge reasoning, point-cloud features, multi-modal cues, spatial knowledge, self-supervised pretraining,

object-centric features, and edge-centric relational reasoning (Zhang et al. 2021; Feng et al. 2023; Wang et al. 2023; Koch et al. 2024a; Heo et al. 2025; Ma et al. 2026). These methods already use geometry; GeoCalib instead targets the post-source reliability gap that a high semantic relation score need not be calibrated to physical validity for the same object pair.

Open-vocabulary 3D graphs. Open-vocabulary 3D perception provides queryable object and region features (Peng et al. 2023; Takmaz et al. 2023); graph systems then connect such features to compact scene representations for querying, planning, navigation, online mapping, and language interaction (Gu et al. 2024; Werby et al. 2024; Maggio et al. 2024). Recent 3D scene graph work extends this line toward open-vocabulary objects, open-set relations, VLM features, online graph generation, and functional relations (Koch et al. 2024b; Chen et al. 2024; Hou et al. 2025; Zhang et al. 2025; Wang et al. 2025). This line raises the reliability requirement: when relation edges are language-facing scene representations, plausible text labels must still refer to object pairs that satisfy the corresponding spatial relation.

Other recent systems further connect VLMs, zero-shot graph construction, graph-mediated LLM reasoning, retrieval augmentation, and hallucination mitigation (Saxena and Chiun 2025; Madhavaram et al. 2026; Yu et al. 2025; Zemskova and Yudin 2025; Ogunleye, Abdelrahman, and Lourentzou 2026). We cite them as trend and boundary evidence, not direct GeoCalib baselines.


Geometry-aware relation evidence. A broad set of 3D reasoning methods already uses geometry: edge features, support/contact patterns, affordance geometry, metric distance, topology, and semantic-geometric fusion (Zhang et al. 2021; Feng et al. 2023; Huang et al. 2025; Shao et al. 2025; Xie, Pagani, and Stricker 2024). Scene graph alignment and registration methods further demonstrate that graph structure and geometric consistency matter when scene graphs are reused for matching, alignment, and downstream geometric tasks (Sarkar et al. 2023; Xie, Pagani, and Stricker 2024).

RelWitness uses visual-geometric witnesses, calibrated witness quality, and witness-consistent decoding for open-vocabulary 3D scene graph generation (Nguyen et al. 2026). Statistical confidence rescoring refines 3D scene graph predictions from multi-view images using semantic masks, neighboring relations, and statistical priors (Yeo, Li, and Lee 2025). GeoCalib therefore does not claim novelty from visual-geometric evidence, calibration, or rescoring alone. It studies a different post-source contract: identity-preserving geometry joins, recall/violation operating points, and controls over existing relation-source outputs.

Reliability and calibration. Standard $R@K$ asks whether the correct predicate appears near the top, not whether the predicted relation is physically valid in the reconstructed scene. Confidence calibration asks whether model scores reflect empirical correctness likelihood (Guo et al. 2017), and selective prediction makes risk and coverage explicit (Geifman and El-Yaniv 2017). GeoCalib instantiates this reliabil-

GeoCalib: calibrating geometric consistency for relation rows

A semantically plausible edge is not necessarily reliable for the same 3D object pair.



Paper-facing names: product + rank-average = soft instantiations; pooled = ablation; strict rule = diagnostic.

Caption guard: scoped relation-reliability framework, not a broad open-vocabulary generation method.

Figure 1: **GeoCalib in one relation-row example.** A semantically plausible relation can be physically inconsistent for the same object pair. GeoCalib separates predicate semantics T , raw same-pair geometry G , and source confidence Z , calibrates $C(T, G)$ without Z , and supports calibrated-product and scale-robust rank-average re-ranking while reporting recall and violation together.

ity view for geometry-checkable 3DSSG families by standardizing prediction rows, joining same-pair 3D evidence, estimating $p^{\text{geom_valid}}$, and reporting recall with violation.

Problem Formulation

Let a relation source produce a set of candidate relation rows for a 3D scene. Each row is written as

$$r_i = (\text{scan}_i, \text{subgraph}_i, s_i, o_i, p_i, a_i, \text{score}_i^{\text{sem}}), \quad (1)$$

where s_i and o_i are subject and object instance identifiers, p_i is the predicate label, a_i is the mapped predicate family, and $\text{score}_i^{\text{sem}}$ is the source semantic confidence. The tuple $(\text{scan}_i, \text{subgraph}_i, s_i, o_i)$ is the identity key used to preserve the object pair across prediction, geometry, and ground-truth joins. We expose three factors:

$$T_i = (p_i, a_i), \quad Z_i = \text{score}_i^{\text{sem}}, \quad (2)$$

where T_i carries predicate/family semantics and Z_i is source confidence.

For each prediction row, we join raw object-pair geometry evidence

$$G_i = G(\text{scan}_i, \text{subgraph}_i, s_i, o_i). \quad (3)$$

The raw evidence is predicate-independent and can include OBB features, point/local contact evidence, distance, vertical ordering, support/contact subtype information, and missing-evidence reason codes. GeoCalib evaluates only geometry-checkable relation families:

$$\mathcal{A} = \{\text{support_contact}, \text{proximity}, \text{relative_vertical}\}. \quad (4)$$

Rows outside these families are preserved in exports but excluded from the measured violation claim.

The geometry verifier assigns a row-level status $v_i \in \{\text{satisfied}, \text{uncertain}, \text{violated}\}$. The calibrator learns a constructed-target compatibility:

$$C_i = P(y_i^{\text{cal}} = 1 \mid T_i, G_i), \quad Z_i \notin \text{inputs}(C_i). \quad (5)$$

Here y^{cal} contains train-dev GT positives and high-margin counterfactual negatives, not direct human validity. We denote family-conditioned and pooled implementations by $p^{\text{geom_valid}, \text{family}}$ and $p^{\text{geom_valid}}$. Predicate-aligned transformations are $T_i \times G_i$ interactions rather than raw geometry; neither Z_i nor source identity enters the calibrator. Final scores use $S_i = F(Z_i, C_i)$, with product and rank-average as soft instantiations rather than proofs or a universal formula.

The main recall metric is exact predicate-label R@K over the fixed in-scope ground-truth denominator. Family grouping defines the geometry-checkable reliability scope and violation analysis; it does not relax predicate-label matching. Violation@K is the fraction of top-K in-scope predicted rows whose joined geometry is judged violated. Top-K is formed per evaluated subgraph/source ranking after the measured-family filter; satisfied, uncertain, and violated rows form the violation denominator, while unsupported or missing-join rows are excluded from that denominator and reported through coverage caveats. Only rows with verification status violated count as violations. We report recall retention and violation reduction together so that violation reduction cannot be interpreted without its recall tradeoff.

Method

GeoCalib is a framework for 3D scene graph relation predictions. It is not a replacement relation predictor and not a standalone hand-coded verifier. The geometric checks define the measurable validity target for families with observable physical constraints; the method contribution is the identity-preserving row contract, factor separation, calibrated compatibility, operating-point analysis, and controls. The design preserves row identity, joins predicate-independent same-pair geometry, excludes source confidence from compatibility calibration, and evaluates recall and violation jointly. Figure 1 presents this path to calibrated operating points.

Relation-Row Standardization

Different relation sources expose predictions in different formats. We convert each source output into a shared row contract:

$$(\text{scan}, \text{subgraph}, s, o, p, a, \text{score}^{\text{sem}}, \text{source}). \quad (6)$$

This row is the unit of comparison across VL-SAT, Open3DSG, controls, and future optional sources. The standardized row keeps the semantic score from the source model but adds the fields needed for identity-preserving geometry joins and denominator-stable evaluation.

Identity-Preserving Geometry Join

Geometry is joined by $(\text{scan}, \text{subgraph}, s, o)$, not by category names or unordered object-pair statistics. This matters because the same predicate can be plausible for many object categories while being false for a particular instance pair. The wrong-pair and shuffled-geometry controls directly test this design choice by breaking or perturbing object-pair identity while preserving parts of the geometry distribution.

Family-Specific Geometry Checks

GeoCalib uses relation-family-specific checks for `support_contact`, `proximity`, and `relative_vertical`. For `support_contact`, the evidence includes point/local contact, vertical support, OBB overlap, and support subtypes. For `proximity`, the evidence includes object distance and pair scale. For `relative_vertical`, the evidence includes centroid and extent ordering along the vertical axis. The verifier returns violated only when the evidence is strong enough; otherwise ambiguous or weak geometry can be marked uncertain. We do not apply this protocol to predicates whose correctness depends primarily on function, social context, affordance, or language pragmatics, because those predicates do not have a stable geometric validity test from the available reconstruction alone.

Factor-Isolated Calibration and Fusion

For each row, T_i denotes predicate/family semantics, G_i raw predicate-independent same-pair geometry, and Z_i source confidence. The frozen family-conditioned calibrator is

$$p_i^{\text{geom_valid, family}} = C_i = C_{a_i}(\Phi(T_i, G_i)), \quad (7)$$

$$= P(y_i^{\text{cal}} = 1 \mid T_i, G_i),$$

Here $\Phi = (\phi_T, \phi_G, \phi_\times)$; $\phi_\times$ contains explicit predicate-aligned transforms such as signed vertical displacement while raw G_i remains predicate-independent. The leakage boundary is

$$Z_i \notin \text{inputs}(C_i). \quad (8)$$

Source confidence and identity are absent from calibrator features and enter only after C_i is fixed. A pooled calibrator with the same boundary is the `pooled_risk` ablation. Both use 2,565 train-dev GT positives and 3,244 high-margin counterfactual negatives, so y^{cal} is constructed rather than human validity. Family maps, rules, counterfactuals, and calibrators are frozen from train-dev artifacts without held-out retuning.

We first use geometry validity as a soft risk term. For the calibrated product, $R_i^{\text{geom}} = -\log p_i^{\text{geom_valid, family}}$ and

$$U_\lambda(i) = \log \text{score}_i^{\text{sem}} - \lambda R_i^{\text{geom}}. \quad (9)$$

The reported product fixes $\lambda = 1$, giving

$$\text{score}_i^{\text{product}} = \text{score}_i^{\text{sem}} \cdot p_i^{\text{geom_valid, family}}. \quad (10)$$

This preserves the source ranking signal while continuously penalizing calibrated inconsistency; λ is not tuned on held-out source results.

The second soft instantiation removes source-scale assumptions. Within each subgraph, let q_i^{sem} and q_i^{geom} be descending percentile ranks of $\text{score}_i^{\text{sem}}$ and $p_i^{\text{geom_valid, family}}$. Rank-average assigns

$$\text{score}_i^{\text{rank}} = \frac{1}{2}(q_i^{\text{sem}} + q_i^{\text{geom}}). \quad (11)$$

The equal weight is parameter-free and fixed before SGFN inference. The product uses calibrated probability magnitudes; rank-average is scale-robust. We treat them as two GeoCalib fusion instantiations rather than claiming one formula is universally optimal. We additionally report `probabilistic_recalibrated` as the pooled calibrated-risk ablation $\text{score}_i^{\text{sem}} \cdot p_i^{\text{geom_valid}}$, fixed-constant Reciprocal Rank Fusion as a strong comparison, and `rule_verified_point_subtype` as a hard zero-violation diagnostic.

Controls

The control conditions test whether the result can be explained by simpler signals. The legacy $p^{\text{geom_valid}}$ -only ranking removes Z_i and is therefore distinct from the pooled calibrated-risk ablation, which still multiplies source confidence by $p^{\text{geom_valid}}$. We call it calibrator-only, not true G -only, because the calibrator retains T_i and $T_i \times G_i$ features. Distance-only ranking tests whether a generic spatial heuristic explains the result. Shuffled-geometry and wrong-pair-geometry controls test whether relation-level object-pair identity is necessary. If these controls perform worse than the calibrated setting, the result supports the framework design rather than a generic “add geometry” explanation.

Experimental Setup

This section fixes scope before results because GeoCalib’s claim depends on denominator transparency as much as metric values. Open3DSG is the main open-vocabulary relation-source case study, while VL-SAT is a controlled reproduced closed-set anchor. Families and denominators are

fixed before source-result reporting, yielding a predeclared physically checkable slice: 157 scans, 548 contexts, 36,808 directed pairs, 957,008 VL-SAT prediction rows, 11,254 ground-truth rows, and 3,972 in-scope GT relations. Table 1 records the denominator and source outputs used by Table 2.

Open3DSG is reported as a Docker-reproduced source-output reliability case study, not an official leaderboard reproduction. The selected checkpoint is the official non-averaged BLIP route `epoch=13-step=13104.ckpt`, chosen by train-dev validation loss before held-out source-result inspection. Its runtime uses filtered train/dev subgraphs, and the main full-validation branch reaches 548/548 contexts with `OPEN3DSG_MIN_VISIBLE_OBJECTS=2` plus relaxed view regeneration for two scans. The 533/548 covered branch and historical old 377/388 vs R2 388/388 branch remain sensitivity evidence only.

Condition names are fixed before reporting. The calibrated product `family_conditional_risk` ranks by $\text{score}^{\text{sem}} \cdot p^{\text{geom_valid, family}}$; rank-average combines within-subgraph semantic and family-geometry percentiles with fixed equal weight. `probabilistic_recalibrated` is the pooled calibrated-risk ablation $\text{score}^{\text{sem}} \cdot p^{\text{geom_valid}}$. RRF is a fixed strong comparison, while the calibrator-only control ranks by $p^{\text{geom_valid}}$ without source confidence. It is not true raw- G -only because the frozen calibrator retains predicate/family and predicate-aligned interaction features.

Calibration inputs are classified as T , raw predicate-independent G , or $T \times G$; Z and source identity are excluded. A separate strict train-only factor audit does not alter the frozen SGFN gate.

SGFN `full_l160` is the prospective exact-label target. Its source test list exactly matches the frozen 157-scan H001 scope. The target, $K=100$, families, denominator, calibrated product, rank-average, RRF, and bootstrap gates were fixed before correct-checkpoint download and inference. Two pre-preference errata corrected the split identity and an initially mistaken 120 checkpoint URL; all versions remain in provenance. SGFN covers all 36,808 nonself directed pairs. Eleven self-supported by GT rows have no model edge, receive no synthesized prediction, and remain in the 3,972-row recall denominator.

Open3DSG and VL-SAT retain the case-study/anchor roles above. The untouched ReplicaSSG/FROSS transfer stress test is negative evidence (supplement), while Qwen-VL remains optional. Broad 3DSSG generation gains are not claimed. All promoted results use Docker; host-only outputs are excluded.

Blinded Codex LLM proxy audit. Following LLM annotator/evaluator and 3D-judge precedent (Gilardi, Alizadeh, and Kubli 2023; Liu et al. 2023; Wu et al. 2024), two locked Codex passes label 488 items from 137 scans as `valid/invalid/ambiguous/unobservable` from blinded relation/pair evidence; pass two does not see pass one before lock. Their shared model family makes this a proxy stability test, not human evidence.




Figure 2: **Recall-violation tradeoff.** Lower Violation@100 and higher R@100 are better. Open3DSG is reported on the full-validation recovery branch; 533/548 and historical 377/388 vs R2 388/388 branches are sensitivity evidence only. The plotted soft operating point is the calibrated product; Table 4 compares fusion forms.

Results and Discussion

Table 2 asks whether calibrated geometry consistency reduces violations without collapsing recall. The Open3DSG source exposes the failure most clearly: semantic-only ranking places many geometrically inconsistent edges at the top of the list, especially for the low- K diagnostic. GeoCalib changes that operating point rather than merely pruning edges. The calibrated product at $K = 100$ improves recall by 8.86 percentage points over the semantic source and reduces violation by 9.01 points; at $K = 5$ and $K = 10$, it also raises recall while moving violation from the high-error semantic regime to the calibrated-risk regime. Pooled risk reduces violations but is weaker than the product on this tradeoff, and the strict-rule row is a zero-violation diagnostic rather than the default method.

VL-SAT provides the controlled anchor. Its semantic ranking is already close to ceiling, so the expected room for recall improvement is small. GeoCalib still reduces violation consistently across the reported K grid and gives a positive R@100 shift. The pooled ablation slightly favors recall at the largest K but leaves more violation, so we report it as a recall-favoring ablation rather than a soft framework instantiation.

A 1,000-resample Docker subgraph bootstrap supports the same reading: Open3DSG has positive recall and negative violation shifts across K , while near-ceiling VL-SAT has smaller stable violation reductions. Main tables retain the standard fixed-source top- K point estimates.

No fusion dominates retrospectively: product has lower V on Open3DSG, while rank fusion is stronger on VL-SAT at $K = 100$. On frozen SGFN, product changes R/V from 0.9235/0.0630 to 0.9416/0.0381 (paired deltas 0.0181 [0.0134,0.0233] and -0.0249 [$-0.0270, -0.0229$]). Rank-average reaches 0.9476/0.0277; versus product its recall CI crosses zero and V CI is negative. This supports multiple fusion forms, not product superiority.

The aggregate effect is not family-uniform. On SGFN, product `support_contact` verifier V increases by 0.00450 (95% CI [0.00370,0.00532]), `proximity` is un-

Table 1: **Shared GeoCalib evaluation scope.** Recall is exact predicate-label recall over the 3,972-row in-scope GT denominator; family grouping is used for reliability and violation reporting, not for relaxing predicate matches.

Item	Count	Role in claim
Validation scans	157	Shared scope; SGFN’s official test list matches it exactly.
Subgraphs / contexts	548	Object-pair identity denominator.
Directed object pairs	36,808	Candidate relation-row universe.
Ground-truth rows	11,254	Full validation relation annotation rows.
In-scope GT rows	3,972	Exact-label denominator for the three measured relation families.
VL-SAT prediction rows	957,008	Controlled reproduced closed-set anchor output.
Open3DSG prediction rows	695,916	Main open-vocabulary source output.
SGFN prediction rows	957,008	Prospectively frozen exact-label confirmation output.
Geometry-scored rows	160,596	Open3DSG geometry-checkable rows joined to calibrated evidence.

Table 2: **Main source results.** Rows use the full official 3,972-row three-family denominator. Open3DSG uses the selected non-averaged checkpoint on the 548/548 recovery branch; 533/548 and historical 377/388 vs R2 388/388 branches are sensitivity evidence only. We report the fixed low-K grid $K \in \{5, 10, 20, 50, 100\}$; $K = 1$ is excluded as a noisy sanity check. Residual calibration risk is reported in the failure analysis rather than hidden in the table. Product is the calibrated-product instantiation; pooled risk is the pooled probabilistic_recalibrated ablation.

Source	Condition	R@K					Violation@K				
		@5	@10	@20	@50	@100	@5	@10	@20	@50	@100
Open3DSG	Semantic	0.0368	0.1002	0.1991	0.4096	0.5161	0.5131	0.3255	0.2088	0.1386	0.1242
Open3DSG	Pooled risk	0.0826	0.1581	0.2603	0.3975	0.5723	0.0628	0.0699	0.0654	0.0606	0.0811
Open3DSG	Strict rule	0.0707	0.1314	0.2422	0.4295	0.5368	0.0000	0.0000	0.0000	0.0000	0.0000
Open3DSG	Product	0.0984	0.1921	0.3291	0.4658	0.6047	0.0420	0.0482	0.0441	0.0286	0.0341
VL-SAT	Semantic	0.4194	0.6322	0.8074	0.9272	0.9635	0.0029	0.0082	0.0142	0.0268	0.0476
VL-SAT	Pooled risk	0.4154	0.6322	0.8107	0.9305	0.9688	0.0015	0.0071	0.0120	0.0229	0.0404
VL-SAT	Strict rule	0.4197	0.6317	0.8074	0.9257	0.9627	0.0000	0.0000	0.0000	0.0000	0.0000
VL-SAT	Product	0.4162	0.6309	0.8087	0.9288	0.9683	0.0011	0.0051	0.0109	0.0206	0.0333

changed, and `relative_vertical` supplies most of the reduction. The same support/contact direction appears in the retrospective sources. We therefore claim aggregate scoped reliability and expose the family decomposition rather than claiming every-family improvement.

Table 3 rules out several simpler alternatives. Calibrator-only is not pooled calibrated re-ranking: it removes source confidence entirely, although it still contains T and $T \times G$. Distance-only is weaker, and shuffled/wrong-pair controls show that object-pair identity is necessary. Thus GeoCalib is not merely ranking by a calibrator, distance, or hard rules. A strict 1,061/117/157 scan train/dev/final audit adds true raw- G , T -only/additive/interaction baselines, and counterfactual controls: correct- T wins 97.44% of 390 vertical GT rows; close-by swap and vertical inverse errors are 0.01542/0.00124; and the correct-wrong-pair gap is +0.42341 over 3,961 rows. This completes the factor audit, but neither solves support/contact nor learns an evidence router.

The GT-based verifier check tests the geometry signal independently: GT positives remain nonviolated at 0.9965, counterfactual negatives are nonsatisfied at 0.9673, and $p^{\text{geom_valid}}$ separates 7,944 rows with AUROC/AUPRC 0.9772/0.9729 and Brier 0.0543. This distinguishes the calibrated reliability signal from a purely hand-written validity script.

Codex proxy passes yield valid/invalid/ambiguous/unobservable counts of 180/185/120/3 and 175/178/132/3. Agreement is 438/488 (89.75%, four-class $\kappa = 0.845$); all 334 jointly binary items match, and every disagreement touches ambiguity. This is blinded same-system stability, not Human V@K or ground truth; earlier structured/visual checks remain supplementary.

Together, the metrics support a reliability layer, not a new generator or broad SOTA claim. Open3DSG comparisons share one checkpoint, row contract, and denominator. The historical 388/388 sensitivity branch changes old 377/388 R@100 by only +0.28 points; it stays outside the main table, with its teardown exit-137 caveat preserved in provenance.

Failure analysis explains why the Open3DSG metrics change. The full-validation branch produces 82,155 diagnostic rows with zero validation errors; in the 36-case queue, 25 cases are demoted by geometry-aware re-ranking. Figure 3 shows distance, vertical-order, and support-plane contradictions that make symbolic edges unreliable for downstream reasoning.

The same analysis exposes residual risk: ten of 36 sampled rule-violated cases still have $p^{\text{geom_valid}} > 0.9$. Thus $p^{\text{geom_valid}}$ is not a hard validity label, and the paper reports both soft fusion instantiations alongside pooled, RRF, and rule-verified variants.

Table 3: **Controls and diagnostics.** Calibrator-only ranks by $p^{\text{geom_valid}}$ without source confidence; pooled risk keeps source confidence and uses the pooled calibrator. The former is not true raw- G -only because its calibrator retains T and $T \times G$. Shuffled and wrong-pair controls break identity-preserving geometry evidence.

Source	Condition	R@100	Violation@100	Interpretation
Open3DSG	Calibrator only	0.5116	0.0865	Removing Z does not recover the main tradeoff.
Open3DSG	Distance-only	0.5038	0.1071	Generic distance is weaker than calibrated geometry risk.
Open3DSG	Shuffled geometry	0.2543	0.1998	Breaking row-level geometry destroys reliability.
Open3DSG	Wrong-pair geometry	0.2331	0.1985	Same-pair identity is necessary.
VL-SAT	Calibrator only	0.5184	0.0711	Removing Z loses substantial recall.
VL-SAT	Distance-only	0.5554	0.0981	Distance is a poor proxy for relation reliability.
VL-SAT	Shuffled geometry	0.9494	0.0588	Violations increase when geometry identity is broken.
VL-SAT	Wrong-pair geometry	0.9529	0.0601	Wrong object-pair evidence is worse than GeoCalib.

Table 4: **Prospective framework confirmation at $K = 100$.** Product and rank-average are GeoCalib soft instantiations; fixed-constant RRF is a strong comparison. SGFN was prospectively frozen and uses the unchanged 3,972-row denominator. V is verifier-derived, not independent human validity.

Source	Condition	R@100	Violation@100	Interpretation
SGFN	Semantic	0.9235	0.0630	Fresh exact-label source baseline.
SGFN	Product	0.9416	0.0381	Frozen joint gate pass vs semantic.
SGFN	Rank-average	0.9476	0.0277	Frozen joint gate pass vs product.
SGFN	RRF	0.9192	0.0284	Lower V but fails the recall guardrail.

Figure 3: Geometry-backed failure mechanism examples

Same locked Open3DSG qualitative rows, now drawn from preprocessed object point clouds.




Figure 3: **Geometry-backed qualitative failures.** Open3DSG examples show proximity, relative-vertical, and support/contact cases where semantic plausibility conflicts with object-pair geometry. The rightmost case illustrates residual calibration risk.

Limitations

The claim covers contact/support, proximity, and vertical ordering, where reconstructed geometry supports falsifiable checks. Functional and pragmatic predicates remain outside scope. Front/behind is frame-ambiguous, and a left/right variant fails its train/dev strict gate; these are reported as negative scope evidence. Exact-label GT-object evaluation also makes Open3DSG relation-reliability evidence rather than full open-vocabulary generation.

Open3DSG uses a selected non-averaged checkpoint and documented 548/548 recovery branch, not an official leaderboard reproduction. The 533/548 route and historical 377/388–388/388 comparison are sensitivity evidence; broad SOTA and arbitrary-source generality remain blocked.

V is verifier-derived; two blinded Codex passes share a model family and lack human alignment, so they are proxy diagnostics. Fusion is source-dependent: support_contact V regresses, while on untouched 11-scene ReplicaSSG/FROSS, product is unchanged at $K = 100$ and rank-average fails the recall guardrail (supplement). Thus we claim no formula, family, or dataset-wide generality.

SGFN preserves two pre-inference errata for split identity and the full_1160 URL. Eleven self-supported by GT rows lack model edges; no predictions are synthesized, and all remain in the 3,972-row recall denominator. The scores remain constructed-target reliability, not human validity; factor controls do not expand SGFN or support a generic interaction claim. Future work requires human alignment, broader relations, and embodied tests.

Conclusion

GeoCalib addresses semantic–geometry mismatch by calibrating $C(T, G)$ without source confidence Z and evaluating $F(Z, C)$ with recall and violation. Across three sources, results support the scoped framework, not formula dominance.

References

Armeni, I.; He, Z.-Y.; Gwak, J.; Zamir, A. R.; Fischer, M.; Malik, J.; and Savarese, S. 2019. 3D Scene Graph: A Structure for Unified Semantics, 3D Space, and Camera. In *Pro-*

ceedings of the *IEEE/CVF International Conference on Computer Vision (ICCV)*, 5664–5673.

Chen, L.; Wang, X.; Lu, J.; Lin, S.; Wang, C.; and He, G. 2024. CLIP-Driven Open-Vocabulary 3D Scene Graph Generation via Cross-Modality Contrastive Learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 27863–27873.

Feng, M.; Hou, H.; Zhang, L.; Wu, Z.; Guo, Y.; and Mian, A. 2023. 3D Spatial Multimodal Knowledge Accumulation for Scene Graph Prediction in Point Cloud. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 9182–9191.

Geifman, Y.; and El-Yaniv, R. 2017. Selective Classification for Deep Neural Networks. In *Advances in Neural Information Processing Systems*, volume 30.

Gilardi, F.; Alizadeh, M.; and Kubli, M. 2023. ChatGPT Outperforms Crowd Workers for Text-Annotation Tasks. *Proceedings of the National Academy of Sciences*, 120(30): e2305016120.

Gu, Q.; Kuwajerwala, A.; Morin, S.; Jatavallabhula, K. M.; Sen, B.; Agarwal, A.; Rivera, C.; Paul, W.; Ellis, K.; Chellappa, R.; Gan, C.; de Melo, C. M.; Tenenbaum, J. B.; Torralba, A.; Shkurti, F.; and Paull, L. 2024. ConceptGraphs: Open-Vocabulary 3D Scene Graphs for Perception and Planning. In *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, 5021–5028.

Guo, C.; Pleiss, G.; Sun, Y.; and Weinberger, K. Q. 2017. On Calibration of Modern Neural Networks. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, 1321–1330.

Heo, K.; Kim, G.; Kim, S.; and Cho, M. 2025. Object-Centric Representation Learning for Enhanced 3D Semantic Scene Graph Prediction. In *Advances in Neural Information Processing Systems*.

Hou, H.-Y.; Lee, C.-Y.; Sonogashira, M.; and Kawanishi, Y. 2025. FROSS: Faster-Than-Real-Time Online 3D Semantic Scene Graph Generation from RGB-D Images. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 28818–28827.

Huang, I.; Bao, Y.; Truong, K.; Zhou, H.; Schmid, C.; Guibas, L.; and Fathi, A. 2025. FirePlace: Geometric Refinements of LLM Common Sense Reasoning for 3D Object Placement. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 13466–13476.

Hughes, N.; Chang, Y.; and Carlone, L. 2022. Hydra: A Real-Time Spatial Perception System for 3D Scene Graph Construction and Optimization. In *Robotics: Science and Systems (RSS)*.

Koch, S.; Hermosilla, P.; Vaskevicius, N.; Colosi, M.; and Ropinski, T. 2024a. SGRec3D: Self-Supervised 3D Scene Graph Learning via Object-Level Scene Reconstruction. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 3404–3414.

Koch, S.; Vaskevicius, N.; Colosi, M.; Hermosilla, P.; and Ropinski, T. 2024b. Open3DSG: Open-Vocabulary 3D Scene

Graphs from Point Clouds with Queryable Objects and Open-Set Relationships. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 14183–14193.

Liu, Y.; Iter, D.; Xu, Y.; Wang, S.; Xu, R.; and Zhu, C. 2023. G-Eval: NLG Evaluation Using GPT-4 with Better Human Alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2511–2522.

Liu, Y.; Mei, H.; Zhan, D.; Zhao, J.; Zhou, D.; Dong, B.; and Yang, X. 2026. View-on-Graph: Zero-Shot 3D Visual Grounding via Vision-Language Reasoning on Scene Graphs. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, 7386–7394.

Ma, Y.; Liu, H.; Guo, Y.; Gevers, T.; and Oswald, M. R. 2026. Edge-Centric Relational Reasoning for 3D Scene Graph Prediction. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, 7847–7855.

Madhavaram, V.; Sengar, V.; De, A.; and Sharma, C. 2026. VIZOR: Viewpoint-Invariant Zero-Shot Scene Graph Generation for 3D Scene Reasoning. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 8584–8595.

Maggio, D.; Chang, Y.; Hughes, N.; Trang, M.; Griffith, J. D.; Dougherty, C.; Cristofalo, E.; Schmid, L.; and Carlone, L. 2024. Clío: Real-Time Task-Driven Open-Set 3D Scene Graphs. *IEEE Robotics and Automation Letters*, 9(10): 8921–8928.

Nguyen, M. A.; Tran, Q. H.; Le, B. N.; Pham, T. K.; and Guang, S. Y. 2026. RelWitness: Open-Vocabulary 3D Scene Graph Generation with Visual-Geometric Relation Witnesses. arXiv:2605.20823.

Ogunleye, M. A.; Abdelrahman, E.; and Lourentzou, I. 2026. 3D-VCD: Hallucination Mitigation in 3D-LLM Embodied Agents through Visual Contrastive Decoding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 40197–40207.

Peng, S.; Genova, K.; Jiang, C. M.; Tagliasacchi, A.; Pollefeys, M.; and Funkhouser, T. 2023. OpenScene: 3D Scene Understanding With Open Vocabularies. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 815–824.

Sarkar, S. D.; Miksik, O.; Pollefeys, M.; Barath, D.; and Armeni, I. 2023. SGAligner: 3D Scene Alignment with Scene Graphs. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 21927–21937.

Saxena, P.; and Chiun, J. 2025. ZING-3D: Zero-shot Incremental 3D Scene Graphs via Vision-Language Models. arXiv:2510.21069.

Shao, Y.; Zhai, W.; Yang, Y.; Luo, H.; Cao, Y.; and Zha, Z.-J. 2025. GREAT: Geometry-Intention Collaborative Inference for Open-Vocabulary 3D Object Affordance Grounding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 17326–17336.

Takmaz, A.; Fedele, E.; Sumner, R. W.; Pollefeys, M.; Tombari, F.; and Engelmann, F. 2023. OpenMask3D: Open-Vocabulary 3D Instance Segmentation. In *Advances in Neu-*

ral Information Processing Systems, volume 36, 68367–68390.

Wald, J.; Dhano, H.; Navab, N.; and Tombari, F. 2020. Learning 3D Semantic Scene Graphs From 3D Indoor Reconstructions. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 3961–3970.

Wang, Z.; Cheng, B.; Zhao, L.; Xu, D.; Tang, Y.; and Sheng, L. 2023. VL-SAT: Visual-Linguistic Semantics Assisted Training for 3D Semantic Scene Graph Prediction in Point Cloud. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 21560–21569.

Wang, Z.; Su, Y.; Li, C.; Wang, D.; Huang, Y.; Li, X.; and Zhao, B. 2025. Open-Vocabulary Octree-Graph for 3D Scene Understanding. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 7037–7047.

Werby, A.; Huang, C.; Büchner, M.; Valada, A.; and Burgard, W. 2024. Hierarchical Open-Vocabulary 3D Scene Graphs for Language-Grounded Robot Navigation. In *Robotics: Science and Systems (RSS)*.

Wu, S.-C.; Wald, J.; Tateno, K.; Navab, N.; and Tombari, F. 2021. SceneGraphFusion: Incremental 3D Scene Graph Prediction From RGB-D Sequences. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 7515–7525.

Wu, T.; Yang, G.; Li, Z.; Zhang, K.; Liu, Z.; Guibas, L.; Lin, D.; and Wetzstein, G. 2024. GPT-4V(ision) Is a Human-Aligned Evaluator for Text-to-3D Generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 22227–22238.

Xie, Y.; Pagani, A.; and Stricker, D. 2024. SG-PGM: Partial Graph Matching Network with Semantic Geometric Fusion for 3D Scene Graph Alignment and Its Downstream Tasks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 28401–28411.

Yeo, Q. X.; Li, Y.; and Lee, G. H. 2025. Statistical Confidence Rescoring for Robust 3D Scene Graph Generation from Multi-View Images. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 24999–25008.

Yu, F.; Deng, Q.; Tang, S.; Li, Y.; and Cheng, L. 2025. Open-World 3D Scene Graph Generation for Retrieval-Augmented Reasoning. Accepted by AAAI 2026, arXiv:2511.05894.

Zemskova, T.; and Yudin, D. 2025. 3DGraphLLM: Combining Semantic Graphs and Large Language Models for 3D Scene Understanding. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 8885–8895.

Zhang, C.; Delitzas, A.; Wang, F.; Zhang, R.; Ji, X.; Pollefeys, M.; and Engelmann, F. 2025. Open-Vocabulary Functional 3D Scene Graphs for Real-World Indoor Spaces. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 19401–19413.

Zhang, C.; Yu, J.; Song, Y.; and Cai, W. 2021. Exploiting Edge-Oriented Reasoning for 3D Point-Based Scene Graph

Analysis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 9705–9715.